

Linear_Regression(1)

November 17, 2023

```
[1]: import pandas as pd
      # importing the pandas library as pd for data manipulation
```

```
[2]: %matplotlib inline
```

```
[3]: # importing the matplotlib library as plt for data visualization
      import matplotlib.pyplot as plt
```

```
[4]: # reading the csv file into a pandas dataframe
      nyc=pd.read_csv('ave_hi_nyc_jan_1895-2018.csv')
```

```
[5]: # displaying the first five rows of the dataframe
      nyc.head()
```

```
[5]:
```

	Date	Value	Anomaly
0	189501	34.2	-3.2
1	189601	34.7	-2.7
2	189701	35.5	-1.9
3	189801	39.6	2.2
4	189901	36.4	-1.0

```
[6]: # displaying the last five rows of the dataframe
      nyc.tail()
```

```
[6]:
```

	Date	Value	Anomaly
119	201401	35.5	-1.9
120	201501	36.1	-1.3
121	201601	40.8	3.4
122	201701	42.8	5.4
123	201801	38.7	1.3

```
[7]: # displaying random 20 rows of the dataframe
      nyc.sample(20)
```

```
[7]:
```

	Date	Value	Anomaly
88	198301	38.0	0.6
21	191601	42.3	4.9

38	193301	45.6	8.2
36	193101	38.2	0.8
61	195601	36.2	-1.2
100	199501	42.6	5.2
4	189901	36.4	-1.0
106	200101	38.8	1.4
108	200301	32.9	-4.5
115	201001	38.1	0.7
94	198901	42.4	5.0
85	198001	38.0	0.6
80	197501	41.7	4.3
42	193701	45.7	8.3
29	192401	39.4	2.0
73	196801	31.7	-5.7
119	201401	35.5	-1.9
47	194201	36.5	-0.9
11	190601	42.3	4.9
66	196101	31.9	-5.5

```
[8]: # declaring the column names of the dataframe
nyc.columns = ['Date', 'Temperature', 'Anomaly']
```

```
[9]: # displaying the first five rows of the dataframe
nyc.head()
```

```
[9]:
```

	Date	Temperature	Anomaly
0	189501	34.2	-3.2
1	189601	34.7	-2.7
2	189701	35.5	-1.9
3	189801	39.6	2.2
4	189901	36.4	-1.0

```
[10]: # Takes the values of the Date column in the nyc dataframe and divides them by
↳100
nyc.Date=nyc.Date.floordiv(100)
```

```
[11]: # displaying the first five rows of the dataframe
nyc.head()
```

```
[11]:
```

	Date	Temperature	Anomaly
0	1895	34.2	-3.2
1	1896	34.7	-2.7
2	1897	35.5	-1.9
3	1898	39.6	2.2
4	1899	36.4	-1.0

```
[12]: # gets the dimensions of the nyc dataframe
nyc.shape
```

```
[12]: (124, 3)
```

```
[13]: # displays the last five rows of the nyc dataframe
nyc.tail()
```

```
[13]:
```

	Date	Temperature	Anomaly
119	2014	35.5	-1.9
120	2015	36.1	-1.3
121	2016	40.8	3.4
122	2017	42.8	5.4
123	2018	38.7	1.3

```
[15]: # importing the train_test_split function from the sklearn.model_selection
      ↳ library
from sklearn.model_selection import train_test_split
```

```
[16]: # This code accesses the Date column of the nyc dataset and returns the shape
      ↳ of the values in that column.
nyc.Date.values.shape
```

```
[16]: (124,)
```

```
[17]: # splitting the nyc dataframe into training and testing sets. The training set
      ↳ will be used to train the model and the testing set will be used to test the
      ↳ model.
X_train,X_test,y_train,y_test=train_test_split(nyc.Date.values.reshape(-1,1),
      ↳ nyc.Temperature.values,random_state=11)
```

```
[18]: # This code accesses the shape of the X_train variable(dimension of the
      ↳ training set)
X_train.shape
```

```
[18]: (93, 1)
```

```
[19]: # This code accesses the shape of the X_test variable(dimension of the testing
      ↳ set)
X_test.shape
```

```
[19]: (31, 1)
```

```
[20]: 93+31
```

```
[20]: 124
```

```
[21]: 93/124*100
```

```
[21]: 75.0
```

```
[22]: #Training the model
```

```
[23]: # importing the LinearRegression class from the sklearn.linear_model library
from sklearn.linear_model import LinearRegression
```

```
[24]: # creating an instance of the LinearRegression class
linear_regression=LinearRegression()
```

```
[25]: # calling the linear_regression object instance.
linear_regression
```

```
[25]: LinearRegression()
```

```
[26]: # training the model using the fit method of the linear_regression object
↳instance
linear_regression.fit(X=X_train, y=y_train)
```

```
[26]: LinearRegression()
```

```
[27]: #Equation M and C
```

```
[28]: # This code accesses the slope of the linear regression model(coefficient of
↳the linear regression model)
linear_regression.coef_
```

```
[28]: array([0.01939167])
```

```
[29]: # intercept of the linear regression equation
linear_regression.intercept_
```

```
[29]: -0.30779820252656265
```

The equation: $Temperature = 0.01939167 * Date - 0.30779820252656265$

```
[30]: # Testing the model
(0.01939167 * 2014) - 0.30779820252656265
```

```
[30]: 38.74702517747344
```

```
[31]: # Testing the Model
```

```
[32]: # This code accesses the predicted values of the linear regression model on
↳x_test
```

```
predicted = linear_regression.predict(X_test)
```

```
[33]: # expected values of the linear regression model on the y_test set
      expected = y_test
```

```
[34]: # This code snippet is iterating over two lists, predicted and expected, and
      ↪printing out each corresponding pair of values.
      for p,e in zip(predicted[:,], expected[:,]):
          print(f'Predicted: {p:.2f}, Expected: {e:.2f}')
```

```
Predicted: 37.86, Expected: 31.70
Predicted: 36.48, Expected: 35.50
Predicted: 37.93, Expected: 40.50
Predicted: 36.61, Expected: 29.80
Predicted: 36.75, Expected: 40.70
Predicted: 38.69, Expected: 34.80
Predicted: 36.44, Expected: 34.20
Predicted: 37.14, Expected: 38.20
Predicted: 37.62, Expected: 36.20
Predicted: 37.53, Expected: 42.50
Predicted: 37.00, Expected: 39.40
Predicted: 38.32, Expected: 40.90
Predicted: 37.20, Expected: 39.80
Predicted: 38.46, Expected: 40.80
Predicted: 36.56, Expected: 37.00
Predicted: 37.25, Expected: 45.70
Predicted: 38.18, Expected: 33.00
Predicted: 37.89, Expected: 29.90
Predicted: 38.15, Expected: 38.00
Predicted: 38.63, Expected: 42.40
Predicted: 38.05, Expected: 32.30
Predicted: 37.02, Expected: 33.80
Predicted: 37.12, Expected: 38.50
Predicted: 37.70, Expected: 37.80
Predicted: 36.73, Expected: 36.10
Predicted: 37.64, Expected: 33.80
Predicted: 37.56, Expected: 42.40
Predicted: 38.11, Expected: 30.60
Predicted: 36.87, Expected: 38.40
Predicted: 36.85, Expected: 42.30
Predicted: 36.94, Expected: 39.70
```

```
[1]: # This code snippet is a for loop that iterates over two lists, predicted and
      ↪expected, and prints the predicted value, expected value, and the difference
      ↪between them (error) for each corresponding pair of values.
      for p,e in zip(predicted[:,], expected[:,]):
          print(f'Predicted: {p:.2f}, Expected: {e:.2f}, Error: {e-p:.2f}')
```

```

-----
NameError                                Traceback (most recent call last)
Cell In[1], line 2
      1 # This code snippet is a for loop that iterates over two lists,
      ↪ predicted and expected, and prints the predicted value, expected value, and
      ↪ the difference between them (error) for each corresponding pair of values. The
      ↪ zip() function is used to iterate over both lists simultaneously. The f-string
      ↪ is used to format the output string with the predicted value, expected value,
      ↪ and error. The :.2f format specifier is used to display the floating-point
      ↪ numbers with two decimal places. This code snippet is a for loop that iterates
      ↪ over two lists, predicted and expected, and prints the predicted value,
      ↪ expected value, and the difference between them (error) for each corresponding
      ↪ pair of values. The zip() function is used to iterate over both lists
      ↪ simultaneously. The f-string is used to format the output string with the
      ↪ predicted value, expected value, and error. The :.2f format specifier is used
      ↪ to display the floating-point numbers with two decimal places.
----> 2 for p,e in zip(predicted[:, :], expected[:, :]):
      3     print(f'Predicted: {p:.2f}, Expected: {e:.2f}, Error: {e-p:.2f}')

NameError: name 'predicted' is not defined

```

```
[36]: # Mean Absolute Error (MAE)
```

```
[37]: # importing the mean_absolute_error function from the sklearn.metrics library.
      ↪ Used to calculate the mean absolute error of the linear regression model
from sklearn.metrics import mean_absolute_error
```

```
[38]: # printing the mean absolute error of the linear regression model
print("MAE", mean_absolute_error(expected, predicted))
```

MAE 3.450753706948741

```
[39]: #RMSE - Root Mean Squared Error
```

```
[40]: # importing the mean_squared_error function from the sklearn.metrics library.
      ↪ Used to calculate the mean squared error of the linear regression model
from sklearn.metrics import mean_squared_error
```

```
[41]: # importing numpy as np for mathematical operations on arrays
import numpy as np
```

```
[42]: # printing the root mean squared error of the linear regression model
print(f'RMSE', np.log(np.sqrt(mean_squared_error(expected, predicted))) )
```

RMSE 1.4256936366057185

```
[43]: # R Squared (R2)
```

```
[44]: # importing the r2_score function from the sklearn.metrics library. Used to
      ↪ calculate the r2 score of the linear regression model
      from sklearn.metrics import r2_score
```

```
[45]: # r2 variable stores the r2 score function of the linear regression model
      r2=r2_score(expected,predicted)
```

```
[46]: r2
```

```
[46]: -0.033370346388810423
```

```
[47]: # predict future model
```

```
[48]: # The code you provided defines a lambda function named predict.
      # The lambda function takes a single argument, x, and returns the result of the
      ↪ linear regression model's coef_ attribute multiplied by x plus the
      ↪ intercept_ attribute.
```

```
predict=(lambda x: linear_regression.coef_ *x + linear_regression.intercept_)
```

```
[49]: # This code snippet calls the predict function and passes in the value 2014.
      predict(2014)
```

```
[49]: array([38.74703181])
```

```
[50]: # This code snippet calls the predict function and passes in the value 2022.
      predict(2022)
```

```
[50]: array([38.9021652])
```

```
[51]: # This code snippet calls the predict function and passes in the value 1800.
      predict(1800)
```

```
[51]: array([34.59721373])
```

```
[52]: # Visualizing the dataset with a Regression Line
```

```
[54]: # importing the seaborn library as sns for data visualization
      import seaborn as sns
```

```
Requirement already satisfied: seaborn in
/home/malims/Zen/myenv/lib/python3.11/site-packages (0.13.0)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from seaborn) (1.25.2)
Requirement already satisfied: pandas>=1.2 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from seaborn) (2.0.3)
```

Requirement already satisfied: matplotlib!=3.6.1,>=3.3 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from seaborn) (3.7.2)

Requirement already satisfied: contourpy>=1.0.1 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.3->seaborn) (1.1.0)

Requirement already satisfied: cycler>=0.10 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.3->seaborn) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.3->seaborn) (4.41.1)

Requirement already satisfied: kiwisolver>=1.0.1 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.3->seaborn) (1.4.4)

Requirement already satisfied: packaging>=20.0 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.3->seaborn) (23.1)

Requirement already satisfied: pillow>=6.2.0 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.3->seaborn) (10.0.0)

Requirement already satisfied: pyparsing<3.1,>=2.3.1 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.3->seaborn) (3.0.9)


Requirement already satisfied: python-dateutil>=2.7 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.3->seaborn) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from pandas>=1.2->seaborn)
(2023.3)

Requirement already satisfied: tzdata>=2022.1 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from pandas>=1.2->seaborn)
(2023.3)

Requirement already satisfied: six>=1.5 in
/home/malims/Zen/myenv/lib/python3.11/site-packages (from python-
dateutil>=2.7->matplotlib!=3.6.1,>=3.3->seaborn) (1.16.0)

```
[55]: # This code snippet creates a scatter plot of the nyc dataframe's Date and
      ↪ Temperature columns.
axes=sns.scatterplot(data=nyc, x='Date',y='Temperature',
                      hue='Temperature', palette='winter', legend=False)
axes.set_ylim(10,70)
x=np.array([min(nyc.Date.values),max(nyc.Date.values)])
y=predict(x)
line=plt.plot(x,y)
```

```
[56]: # The x variable stores the minimum and maximum values of the nyc dataframe's
      ↪ Date column.
```

```
x
```

```
[56]: array([1895, 2018])
```

```
[57]: # The y variable stores the predicted values of the linear regression model
      ↪ based on the x variable values.
```

```
y
```

```
[57]: array([36.43942269, 38.82459851])
```

```
[ ]:
```