```
In [ ]: import pandas as pd
In [ ]: %matplotlib inline
In [ ]: # importing the matplotlib library as plt for data visualization
        import matplotlib.pyplot as plt
In [ ]: # reading the csv file into a pandas dataframe
        nyc=pd.read_csv('ave_hi_nyc_jan_1895-2018.csv')
In [ ]: # displaying the first five rows of the dataframe
        nyc.head()
Out[]:
             Date Value Anomaly
         0 189501
                    34.2
                             -3.2
        1 189601
                    34.7
                             -2.7
         2 189701
                    35.5
                             -1.9
        3 189801
                    39.6
                              2.2
         4 189901
                    36.4
                             -1.0
In [ ]: # displaying the last five rows of the dataframe
        nyc.tail()
Out[]:
               Date Value Anomaly
         119 201401
                      35.5
                               -1.9
         120 201501
                      36.1
                               -1.3
         121 201601
                      40.8
                                3.4
         122 201701
                                5.4
                      42.8
         123 201801
                      38.7
                                1.3
In [ ]: # displaying random 20 rows of the dataframe
        nyc.sample(20)
```

Out[ ]:		Date	Value	Anomaly
	90	198501	33.0	-4.4
	104	199901	40.8	3.4
	1	189601	34.7	-2.7
	103	199801	45.8	8.4
	82	197701	26.1	-11.3
	51	194601	40.2	2.8
	40	193501	35.0	-2.4
	74	196901	35.7	-1.7
	107	200201	45.4	8.0
	91	198601	39.3	1.9
	19	191401	36.9	-0.5
	78	197301	40.6	3.2
	77	197201	40.5	3.1
	100	199501	42.6	5.2
	113	200801	42.4	5.0
	92	198701	35.9	-1.5
	48	194301	36.1	-1.3
	35	193001	38.5	1.1
	41	193601	34.3	-3.1
	83	197801	32.3	-5.1

```
In [ ]: # declaring the column names of the dataframe
nyc.columns =['Date','Temperature','Anomaly']
```

In [ ]: # displaying the first five rows of the dataframe
 nyc.head()

Out[ ]:		Date	Temperature	Anomaly
	0	189501	34.2	-3.2
	1	189601	34.7	-2.7
	2	189701	35.5	-1.9
	3	189801	39.6	2.2
	4	189901	36.4	-1.0

In [ ]: # displaying the first five rows of the dataframe
 nyc.head()

```
Out[ ]:
           Date Temperature Anomaly
         0 1895
                        34.2
                                 -3.2
         1 1896
                        34.7
                                 -2.7
         2 1897
                        35.5
                                 -1.9
          1898
                        39.6
                                  2.2
           1899
                        36.4
                                 -1.0
In [ ]: # gets the dimensions of the nyc dataframe
        nyc.shape
Out[]: (124, 3)
In [ ]: # displays the last five rows of the nyc dataframe
        nyc.tail()
Out[]:
             Date Temperature Anomaly
                                   -1.9
         119 2014
                          35.5
         120 2015
                          36.1
                                   -1.3
         121 2016
                          40.8
                                    3.4
         122 2017
                          42.8
                                    5.4
         123 2018
                          38.7
                                    1.3
In [ ]: # importing the train test split function from the sklearn.model selection
        from sklearn.model selection import train test split
In [ ]: nyc.Date.values.shape
Out[]: (124,)
In [ ]: # splitting the nyc dataframe into training and testing sets. The training
        X_train, X_test, y_train, y_test=train_test_split(nyc.Date.values.reshape(-1
In [ ]: # This code accesses the shape of the X_train variable(dimension of the t
        X_train.shape
Out[]: (93, 1)
In [ ]: # This code accesses the shape of the X_test variable(dimension of the te
        X_test.shape
Out[]: (31, 1)
In [ ]:
        93+31
Out[]: 124
In [ ]:
        93/124*100
Out[]: 75.0
```

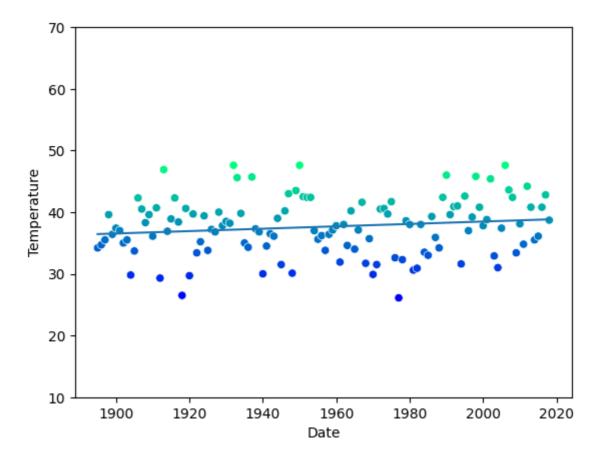
```
In [ ]: #Training the model
In [ ]: # importing the LinearRegression class from the sklearn.linear model libr
        from sklearn.linear model import LinearRegression
In [ ]: # creating an instance of the LinearRegression class
        linear regression=LinearRegression()
In [ ]: # calling the linear regression object instance.
        linear regression
Out[]: ▼ LinearRegression
        LinearRegression()
In [ ]: # training the model using the fit method of the linear regression object
        linear_regression.fit(X=X_train, y=y_train)
Out[]: ▼ LinearRegression
        LinearRegression()
In [ ]: #Equation M and C
In [ ]: # This code accesses the slope of the linear regression model(coefficient
        linear regression.coef
Out[]: array([0.01939167])
In [ ]: # intercept of the linear regression equation
        linear regression.intercept
Out[]: -0.30779820252656265
        The equation: Temperature = 0.01939167 * Date - 0.30779820252656265
In [ ]: # Testing the model
        (0.01939167 * 2014) - 0.30779820252656265
Out[]: 38.74702517747344
In [ ]: # Testing the Model
In [ ]: # This code accesses the predicted values of the linear regression model
        predicted = linear_regression.predict(X_test)
In [ ]: # expected values of the linear regression model on the y_test set
        expected = y_test
In [ ]: # This code snippet is iterating over two lists, predicted and expected,
        for p,e in zip(predicted[::], expected[::]):
          print(f'Predicted: {p:.2f}, Expected: {e:.2f}')
```

```
Predicted: 37.86, Expected: 31.70
       Predicted: 36.48, Expected: 35.50
       Predicted: 37.93, Expected: 40.50
       Predicted: 36.61, Expected: 29.80
       Predicted: 36.75, Expected: 40.70
       Predicted: 38.69, Expected: 34.80
       Predicted: 36.44, Expected: 34.20
       Predicted: 37.14, Expected: 38.20
       Predicted: 37.62, Expected: 36.20
       Predicted: 37.53, Expected: 42.50
       Predicted: 37.00, Expected: 39.40
       Predicted: 38.32, Expected: 40.90
       Predicted: 37.20, Expected: 39.80
       Predicted: 38.46, Expected: 40.80
       Predicted: 36.56, Expected: 37.00
       Predicted: 37.25, Expected: 45.70
       Predicted: 38.18, Expected: 33.00
       Predicted: 37.89, Expected: 29.90
       Predicted: 38.15, Expected: 38.00
       Predicted: 38.63, Expected: 42.40
       Predicted: 38.05, Expected: 32.30
       Predicted: 37.02, Expected: 33.80
       Predicted: 37.12, Expected: 38.50
       Predicted: 37.70, Expected: 37.80
       Predicted: 36.73, Expected: 36.10
       Predicted: 37.64, Expected: 33.80
       Predicted: 37.56, Expected: 42.40
       Predicted: 38.11, Expected: 30.60
       Predicted: 36.87, Expected: 38.40
       Predicted: 36.85, Expected: 42.30
       Predicted: 36.94, Expected: 39.70
In [ ]: # This code snippet is a for loop that iterates over two lists, predicted
```

```
for p,e in zip(predicted[::], expected[::]):
    print(f'Predicted: {p:.2f}, Expected: {e:.2f}, Error: {e-p:.2f}')
```

```
Predicted: 37.86, Expected: 31.70, Error: -6.16
       Predicted: 36.48, Expected: 35.50, Error: -0.98
       Predicted: 37.93, Expected: 40.50, Error: 2.57
       Predicted: 36.61, Expected: 29.80, Error: -6.81
       Predicted: 36.75, Expected: 40.70, Error: 3.95
       Predicted: 38.69, Expected: 34.80, Error: -3.89
       Predicted: 36.44, Expected: 34.20, Error: -2.24
       Predicted: 37.14, Expected: 38.20, Error: 1.06
       Predicted: 37.62, Expected: 36.20, Error: -1.42
       Predicted: 37.53, Expected: 42.50, Error: 4.97
       Predicted: 37.00, Expected: 39.40, Error: 2.40
       Predicted: 38.32, Expected: 40.90, Error: 2.58
       Predicted: 37.20, Expected: 39.80, Error: 2.60
       Predicted: 38.46, Expected: 40.80, Error: 2.34
       Predicted: 36.56, Expected: 37.00, Error: 0.44
       Predicted: 37.25, Expected: 45.70, Error: 8.45
       Predicted: 38.18, Expected: 33.00, Error: -5.18
       Predicted: 37.89, Expected: 29.90, Error: -7.99
       Predicted: 38.15, Expected: 38.00, Error: -0.15
       Predicted: 38.63, Expected: 42.40, Error: 3.77
       Predicted: 38.05, Expected: 32.30, Error: -5.75
       Predicted: 37.02, Expected: 33.80, Error: -3.22
       Predicted: 37.12, Expected: 38.50, Error: 1.38
       Predicted: 37.70, Expected: 37.80, Error: 0.10
       Predicted: 36.73, Expected: 36.10, Error: -0.63
       Predicted: 37.64, Expected: 33.80, Error: -3.84
       Predicted: 37.56, Expected: 42.40, Error: 4.84
       Predicted: 38.11, Expected: 30.60, Error: -7.51
       Predicted: 36.87, Expected: 38.40, Error: 1.53
       Predicted: 36.85, Expected: 42.30, Error: 5.45
       Predicted: 36.94, Expected: 39.70, Error: 2.76
In [ ]: # Mean Absolute Error (MAE)
In [ ]: # importing the mean absolute error function from the sklearn.metrics lib
        from sklearn.metrics import mean absolute error
In []: # printing the mean absolute error of the linear regression model
        print("MAE", mean absolute error(expected, predicted))
       MAE 3.450753706948741
In [ ]: #RMSE - Root Mean Squared Error
In [ ]: # importing the mean squared error function from the sklearn.metrics libr
        from sklearn.metrics import mean squared error
In [ ]: # importing numpy as np for mathematical operations on arrays
        import numpy as np
In [ ]: # printing the root mean squared error of the linear regression model
        print(f'RMSE',np.log(np.sqrt(mean squared error(expected, predicted)) ))
       RMSE 1.4256936366057185
In [ ]: # R Squared (R2)
```

```
In [ ]: # importing the r2 score function from the sklearn.metrics library. Used
        from sklearn.metrics import r2 score
In [ ]: # r2 variable stores the r2 score function of the linear regression model
        r2=r2 score(expected, predicted)
In [ ]:
Out[]: -0.033370346388810423
In [ ]: # predict future model
In [ ]: # The code you provided defines a lambda function named predict.
        # The lambda function takes a single argument, x, and returns the result
        predict=(lambda x: linear regression.coef *x + linear regression.interce
In [ ]: # This code snippet calls the predict function and passes in the value 20
        predict(2014)
Out[]: array([38.74703181])
In [ ]: # This code snippet calls the predict function and passes in the value 20
        predict(2022)
Out[]: array([38.9021652])
In [ ]: # This code snippet calls the predict function and passes in the value 18
        predict(1800)
Out[]: array([34.59721373])
In [ ]: # Visualizing the dataset with a Regression Line
In [ ]: # importing the seaborn library as sns for data visualization
        import seaborn as sns
In [ ]: # This code snippet creates a scatter plot of the nyc dataframe's Date an
        axes=sns.scatterplot(data=nyc, x='Date',y='Temperature',
                             hue='Temperature', palette='winter', legend=False)
        axes.set ylim(10,70)
        x=np.array([min(nyc.Date.values), max(nyc.Date.values)])
        y=predict(x)
        line=plt.plot(x,y)
```



In [ ]: # The x variable stores the minimum and maximum values of the nyc datafra  $\mathbf{x}$ 

Out[]: array([1895, 2018])

In [ ]: # The y variable stores the predicted values of the linear regression mod
y

Out[]: array([36.43942269, 38.82459851])

In []: